Imperial College London
Department of Earth Science and Engineering
MSc Environmental Data Science and Machine Learning

Comparison of Direct and Cascaded High-Resolution Spatial Downscaling of Wildfire Data

Author:
Tsuyu Huang
tsuyu.huang25@imperial.ac.uk

Supervisors:
Jorge Veiras (jorge@artiocarbon.com)
Prof. Anna Korre (a.korre@imperial.ac.uk)

Code Repository:
https://github.com/ese-ada-lovelace-2025/irp-th25

Company:
Artio
www.artiocarbon.com

August 25, 2026

Abstract

With climate change, high-resolution wildfire data is increasingly vital for risk assessment. This project investigates deep learning for downscaling GOES imagery (2 km) to Landsat-resolution fire masks (30 m), incorporating topographic and land cover covariates to address extreme class imbalance and 66× resolution gap. We evaluate whether cascaded pipeline (GOES → VIIRS 375 m → Landsat) outperforms direct pipeline (GOES → Landsat), comparing a Hybrid Attention U-Net (HAUNet) against an Autoencoder baseline. All configurations capture genuine fire signals, achieving up to 30× higher AUPRC than prevalence-rate baselines. Direct downscaling yields higher overall confidence ranking across both architectures (e.g. HAUNet AUPRC 0.0188 direct vs. 0.0142 cascaded), whereas cascading improves thresholded spatial overlap metrics, most prominently for HAUNet (IoU 0.0089 vs. 0.0044 direct). HAUNet outperforms Autoencoder across both pipelines, and feature ablations confirm elevation covariates contribute to performance. Results indicate the models are better suited for coarse regional risk flagging than precise fire-area prediction, with direct relevance to practical risk applications where under-prediction is costlier than over-prediction.

Declaration of Generative AI Tools

The content and critical analysis presented in this report were written by the student without generative AI.

Generative AI models (Google Gemini 3.6, Anthropic Claude Sonnet 5) were used as developer tools during software implementation. This usage was restricted to:

- **Code Generation & Modularisation:** Drafting code, modularising scripts, and creating visualisation tools.
- **Debugging & Cloud Integration:** Troubleshooting execution errors, optimisation, and assisting with using Google Cloud Services.
- **Implementation Ideation:** Exploring architectural configurations and workflow design possibilities.

Generated codes were manually reviewed, modified, and integrated by the author. References detailing AI usage is in the project repository at `references.md`.

Contents

Declaration of Generative AI Tools **1**

1 Introduction **3**

2 Current State of Remote Sensing Downscaling **3**

3 Data and Materials **4**

4 Methods **5**
4.1 HAUNet (Hybrid Attention U-Net) 5
4.1.1 Hybrid Attention Mechanisms 5
4.1.2 Full-Resolution Covariate Pyramid Branch 6
4.1.3 Multi-Scale Refinement and Feature Fusion 6
4.1.4 Cross-scale Interaction Module (CIM) 7
4.1.5 Architecture Selection Rationale 7
4.2 Autoencoder Baseline 8
4.3 Loss Functions 9
4.3.1 Focal Dice Loss 9
4.4 Training Setup 9
4.5 Computational Environment 10
4.6 Experimental Design 10
4.7 Evaluation Metrics 10

5 Results **13**

6 Discussions **13**
6.1 Overall Performance 13
6.2 Over-Prediction, Thresholding, and Practical Implications 15
6.3 Cascaded vs. Direct Pipeline 15
6.4 HAUNet vs. Autoencoder Architecture Comparison 16
6.5 Effect of Target Dilation 16
6.6 Feature Ablations 17

7 Conclusions **17**

A Appendix **18**
A.1 Focal Dice Loss Formulation 18
A.2 Input GOES and Covariate channels 18
A.3 Prediction Plots 18

1 Introduction

In recent years, rising popularity in Net-Zero initiatives has driven carbon investors to fund reduction projects, such as reforestation (Lambert and Turner, 2024). However, these projects are at risk from wildfires releasing stored carbon back into the atmosphere, a threat escalating in frequency and severity due to climate change (Jones et al., 2020). This causes economic losses for the insurance industry. e.g. The January 2025 Los Angeles County wildfires set an all-time record, causing $40 billion in insured losses (Dwyer and Nam, 2026). In carbon accounting, underestimating fire area leaves projects underinsured, whereas overestimation is conservative (Grassi et al., 2008). This project, conducted with Artio, aims to provide higher-resolution fire data to give better confidence and risk assessment for insurance industry.

Accurate risk assessment requires wildfire data that is both high-spatial and high-temporal resolution. Geostationary (GEO) satellites provide high-temporal frequency but coarse spatial resolution, while Low Earth Orbit (LEO) satellites offer high spatial detail at low temporal frequency. Existing Deep Learning (DL) downscaling methods bridge this trade-off, but the resolution gap between GEO (2 km) and LEO (30 m) is extreme: a single 2 km GOES pixel covers roughly $4,444$ Landsat 30 m pixels. This is an ill-posed, one-to-many super-resolution task prone to physically implausible outputs (Z. Wang et al., 2020).

To mitigate this limitation, this project integrates additional physical information, such as high-resolution topography and landcover to guide downscaling. The project aims to investigate whether cascaded downscaling (2 km to 375 m to 30 m) produces better results compared to direct downscaling (2 km to 30 m). By breaking up the downscaling process, the model trains on a smaller resolution gap, aiming to mitigate error caused by large resolution gaps. This project will compare cascaded and direct pipeline performance of downscaling architectures (hybrid attention U-Net and Autoencoder).

2 Current State of Remote Sensing Downscaling

Traditional downscaling methods often fail to capture the complex non-linear patterns in remote sensing data. DL methods have shown to outperform these approaches (Sdraka et al., 2022). CNN and U-Net architectures have been widely used for their ability to capture spatial features. GAN-based methods have been used where sharpness is prioritised over pixel-wise accuracy. Recently, transformer-based models use long-range attention to capture dependencies beyond a convolution's receptive field (J. Wang et al., 2025), and diffusion-based approaches have showed strong generative ability (Y. Wang et al., 2023).

However, many architectures are optimised for texture-rich, natural images, and not designed for extreme resolution-gaps ($66.7\times$ in this project). They may perform poorly on active wildfire data, where the target signal is extremely sparse and localised. In risk assessment, fire mask is preferred for direct analysis. This creates two challenges: severe class imbalance that biases standard reconstruction loss functions (Qi et al., 2026); and an ill-posed one-to-many mapping (Z. Wang et al., 2020). Architectures are needed that can preserve fire signal against a dominant background and recover coherent fire-front structure.

Closest precedents to this task are McCarthy et al., 2021 and Badhan et al., 2024: trained a U-Net and autoencoder respectively to downscale GOES observations to VIIRS-equivalent 375 m fire mask. Both stop at 375 m resolution, leaving the challenge of downscaling to 30 m resolution unaddressed.

While wildfires are influenced by factors like terrain and landcover (Sullivan, 2017), generic super-resolution models omit this information. Providing high-resolution covariates provides context and constraints for the model, which may help mitigate the ill-posed nature of the downscaling problem.

Models downscaling across large resolution-gaps using coarse thermal imagery alone are

constrained by satellite physical detection limits. Validation of GOES-16 against 30 m Landsat-8 data demonstrates that the active fire detection probability reaches 95% only when coarse 2 km area contains at least 114 active Landsat sub-pixels (F. Li et al., 2020). Detection probability drops below this threshold, coarse channels frequently carrying zero detectable signal for smaller fires. The equivalent 95% detection for VIIRS 375 m requires only 11 active sub-pixels, reducing signal threshold. This motivates staged downscaling: training on smaller resolution-gap, and intermediate VIIRS stage operating at lower detection threshold, meaning the input carries genuine signal where the direct downscaling may not.

Existing literature is dominated by direct downscaling across the full resolution gap. Comparisons of multi-stage pipelines also remain scarce. X. Li et al., 2025 demonstrated that stepwise downscaling of ERA5-Land air temperature (10.5 km $\rightarrow$ 30 m) outperformed direct downscaling. However, they evaluated meteorological variables using tree-based ensembles. To date, no prior work researches cascaded DL downscaling, particularly for wildfire applications. Nor has existing work explored the trade-off between reduced per-stage resolution gap and the use of independent intermediate satellite ground-truths across stages.

Main project objectives are:

1. Develop and evaluate a downscaling framework tailored for active wildfires at 30 m resolution, beyond 375 m, that exceeds no-skill prevalence-rate baseline under the extreme class imbalance and resolution gap in this task.

2. Quantify whether attention-based architecture (HAUNet) has better performance over simpler Autoencoder baseline.

3. Determine whether cascaded pipeline with intermediate VIIRS stage produces better downscaling performance than direct downscaling, addressing the advantages/trade-off between reduced per-stage resolution gap that remains unexplored in literature.

3 Data and Materials

Table 1 shows the data sources for this project. The specifications for the VIIRS, GOES, and Landsat data are outlined in Table 2.

The GOES-16/17 data were obtained via public Google Cloud Platform (GCP) buckets (NOAA, 2017, 2018). DEM, and ESA WorldCover are obtained via Google Earth Engine (GEE) Python API (Gorelick et al., 2017) and exported directly to Google Cloud Storage (GCS). The Landsat-8 fire masks are derived by de Almeida Pereira et al., 2021 and obtained from the author's public dataset (via Google Drive) and re-uploaded to GCS. VIIRS fire detections are obtained from NASA FIRMS archive. Manifests, split assignments and model checkpoints are also stored in GCS.

There are $12,504$ and $1,202$ total VIIRS and Landsat patches respectively. The asymmetry in sample counts is due to different sources: Landsat from pre-existing external dataset, VIIRS generated directly from a larger available archive. This gives Stage 1 (GOES $\rightarrow$ VIIRS) of the cascaded pipeline a larger, more reliable training set, since Stage 1 errors propagate directly into Stage 2's input. But the final 30m downscaling is trained on the same Landsat dataset.

The data was processed so all regions of interest (ROIs) are dynamically reprojected to their localised UTM zone based on centroid coordinates. Detections from the NASA VIIRS S-NPP archive were filtered for "nominal" and "high" confidence points. Detections are clustered into 7680×7680 m grid cells within a 15-minute window of peak activity. VIIRS grid cell pixels record the peak Fire Radiative Power (FRP). Due to extreme class imbalance of the target fire datasets (mean fraction of positive target pixels are 0.85% for VIIRS, 0.06% for Landsat), stratified sampling was applied across temporal extent and fire pixel count bins during data splitting. The closest temporal GOES-16/17 scan within a $\pm$60-minute guard window was queried.

GOES-16/17 channels 7 (3.9 μm, shortwave IR) and 14 (11.2 μm, longwave IR) are used as they correspond to the spectral bands for fire-detection for their sensitivity to fire radiance. Similar spectral logic applied in VIIRS's active-fire algorithm, which similarly relies on a shortwave (I4, 3.74 μm) and longwave (I5, 11.45 μm) channel pair. (Schmidt et al., 2013)

In addition to 3.9 μm (Ch. 7) and 11.2 μm (Ch. 14) GOES brightness temperatures, a Thermal Fire Index (TFI) is computed as additional input channel:

$$TFI = Ch.7 - Ch.14 \tag{1}$$

Explicitly providing TFI reduces the burden on early network layers to discover a known, physically meaningful relationship, aiming to save model capacity and improving convergence.

Table 1: Data Sources and Specifications

Data Source	**Res.**	**Bands Used**	**Role / Context**
Landsat-8 Fire Masks	30 m	Active Fire Mask de Almeida Pereira et al., 2021	Target / ground truth; binary mask (256×256 px)
VIIRS S-NPP	375 m	I4 (3.74 μm, Active Fire)	Intermediate target / cascaded-pipeline input (FRP, NASA FIRMS)
GOES-16/17 (ABI)	2000 m	Ch. 7 (3.9 μm), Ch. 14 (11.2 μm)	Coarse fire input; brightness temperature, high temporal frequency
Copernicus DEM	30 m	DEM	Static covariate; terrain effects on fire spread
ESA WorldCover	10 m	Landcover (10 classes)	Static covariate; land cover map

4 Methods

4.1 HAUNet (Hybrid Attention U-Net)

Among CNN-based approaches, HAUNet (J. Wang et al., 2023) uses hybrid attention to capture both local fine detail and broader contextual structure across multiple scales.

4.1.1 Hybrid Attention Mechanisms

- **Channel Attention Module (CEM):** Computes channel dependencies with $O(C^2)$ complexity. Given feature map $\mathbf{X} \in \mathbb{R}^{C \times H \times W}$, queries $\mathbf{Q}$, keys $\mathbf{K}$, and values $\mathbf{V}$ are generated via 1×1 convolutions followed by 3×3 depthwise convolutions. The channel attention

Table 2: Data Pipeline Configuration Specifications

Parameter	Value	Description
VIIRS Geographic Bounds	$[-170.0, -30.0]$	Longitude limits for the Americas
VIIRS Date Range	2020-01-01 – 2020-12-31	Temporal query window for VIIRS detections
Landsat Geographic Bounds	N. & S. America	Locations of interest for GOES detections
Landsat Date Range	Aug–Sep 2020	Temporal window for Landsat detections de Almeida Pereira et al., 2021
Min. Landsat Active Pixels	10	Minimum active fire pixels to retain a patch
GOES Temporal Guard	60 min	Max. gap between target time and GOES scan
GOES Native Scale	2000 m	Native GOES ABI grid spacing
Landsat Grid Scale	30 m	Resolution, Landsat branch
VIIRS Grid Scale	375 m	Resolution, VIIRS active fire detections
Landsat Dimensions	256×256 px	Target patch size
VIIRS Dimensions	20×20 px	Target size (7680/375)
GOES Dimensions	4×4 px	Target size (7680/2000)

map $\mathbf{A}_c$ is derived as:

$$\mathbf{A}_c = \text{Softmax}\left(\text{ReLU}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\alpha}\right)\right)\mathbf{V} \tag{2}$$

where α is a learnable scaling parameter.

- **Spatial-Channel Attention (S_CEM):** Parallel spatial attention branch operating along transposed dimensions $\mathbf{Q}_s, \mathbf{K}_s, \mathbf{V}_s \in \mathbb{R}^{B\times N_{head}\times HW\times C}$. The full spatial attention step scales quadratically with spatial resolution as $O((HW)^2)$.

4.1.2 Full-Resolution Covariate Pyramid Branch

The original J. Wang et al., 2023 HAUNet applies full spatial attention in S_CEM with $O((HW)^2)$ memory footprint that causes out-of-memory errors on full-resolution inputs. Hence inputs (256×256) are downsampled by factor $D = 4$ (i.e. $256/4 = 64$) before passing to the model.

The parallel Covariate Pyramid Encoder was introduced as modification to the original HAUNet because downsampling input covariates discards the high-resolution detail at 30 m scale. It processes covariates at full-resolution in a separate branch, then progressively fusing into the main decoder path. This branch uses Windowed Spatial-Channel Attention Blocks (Windowed-S_CEM) (Liu et al., 2021) for computational efficiency.

For a feature map partitioned into non-overlapping local windows of size $P \times P$ ($P = 8$), Windowed-S_CEM computes spatial attention independently within each window. The window is shifted sequentially with size $P/2$. This keeps spatial attention complexity to $O(P^4)$ per window, independent of overall spatial input dimensions.

4.1.3 Multi-Scale Refinement and Feature Fusion

The main encoder-decoder path produces low-resolution feature representations (64×64). The final refinement stage produces the final high-resolution output. At each refinement stage, low-resolution features are upsampled (bilinear interpolation, 3×3 convolution), concatenated with matching-scale Covariate Pyramid Encoder features, and fused via a 3×3 convolution back to the working channel dimension.

The final layer of HAUNet refinement network outputs prediction at the resolution of the full covariate grid (256×256). Operating the decoder at full resolution preserves fine-grained spatial

feature representations across skip connections and covariate channels. This is same for Stage 1 (GOES → VIIRS) model output for cascaded pipeline. Rather than constraining the output's resolution to match the coarse VIIRS target grid (20×20), HAUNet outputs high-resolution spatial feature maps. During metric evaluatation of the Stage 1 model, the outputs are pooled to VIIRS resolution to match VIIRS target.

4.1.4 Cross-scale Interaction Module (CIM)

HAUNet replaces standard skip-connections with Channel Interaction Module (CIM) to resolve semantic mismatches across encoder stages. Encoder features (F_1, F_2, F_3) are upsampled to equal dimensions T_1, T_2, T_3 and concatenated into a multiscale token $T_\Sigma \in \mathbb{R}^{H \times W \times 3C}$. CIM computes channel cross-attention $(Q_i K^T)$ using T_i as Queries and T_Σ as Keys/Values, letting each scale aggregate cross-level context before element-wise addition into decoder blocks.

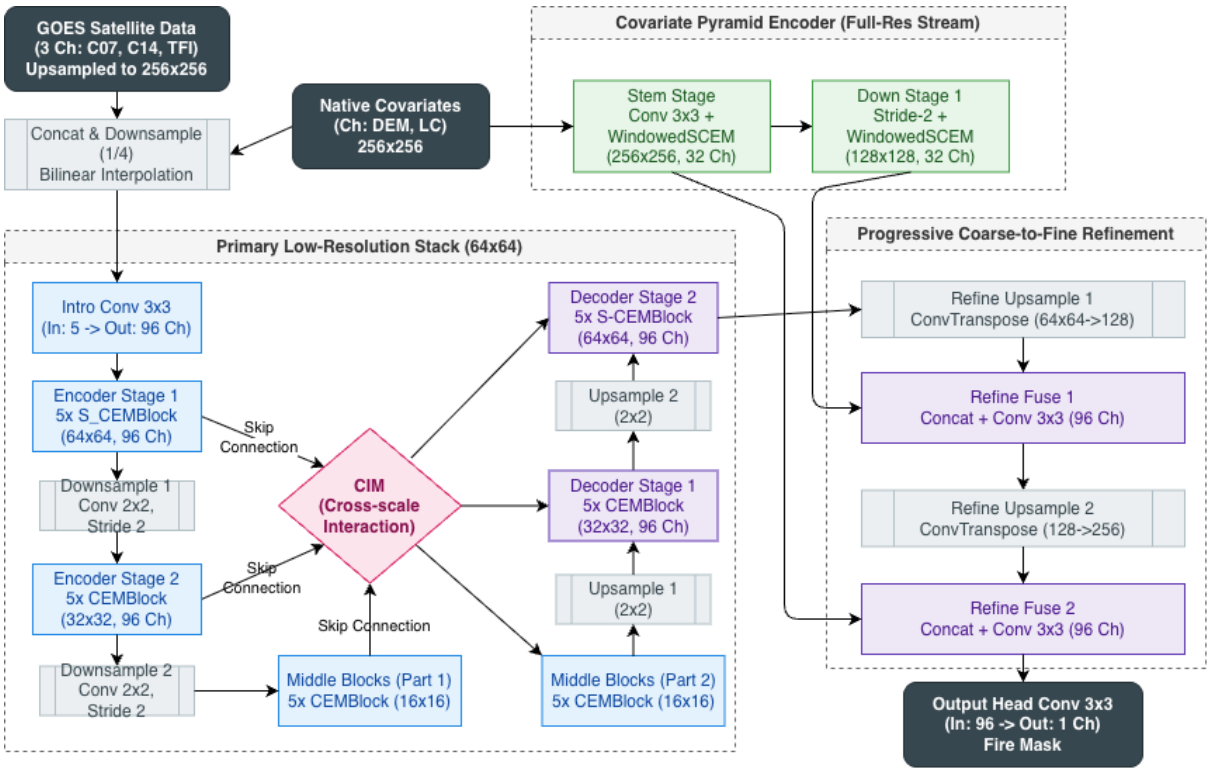


Figure 1: Architecture overview of the wildfire-adapted HAUNet. The full-resolution stream processes native covariates, which are progressively fused via refinement blocks with the low-resolution stack after CIM.

4.1.5 Architecture Selection Rationale

HAUNet was selected as one of the architectures for the following reasons:

1. Dual attention: channel attention to capture global information across feature channels and spatial attention to capture spatial relationships. This is useful for wildfire data, where the fire signal is sparse and localised, but the physical context (e.g. terrain, landcover) is important for accurate downscaling.

2. CIM: standard skip-connections concatenate features at matching spatial resolution, but this assume they are semantically comparable. CIM addresses this by explicitly bridging semantic and resolution gaps between encoder outputs at different scales prior to fusion.

This helps to better learn representation at different scales, which is useful for super-resolution tasks like this (downscaling by 66.7×).

3. Both the autoencoder baseline (4.2) and HAUNet share the similar underlying encoder-decoder-with-skip-connections structure. This allows the comparison between the two models to isolate the contribution of HAUNet's attention mechanisms and skip-connection design.

4.2 Autoencoder Baseline

The baseline architecture follows the autoencoder proposed by Badhan et al., 2024 to downscale from GOES to VIIRS resolution. It is adapted to this project to concatenate the covariates with the fire data inputs, and outputs are fire mask segmentation rather than continuous BT values. The autoencoder follows a standard U-Net-style structure, full architecture is shown in Figure 2.

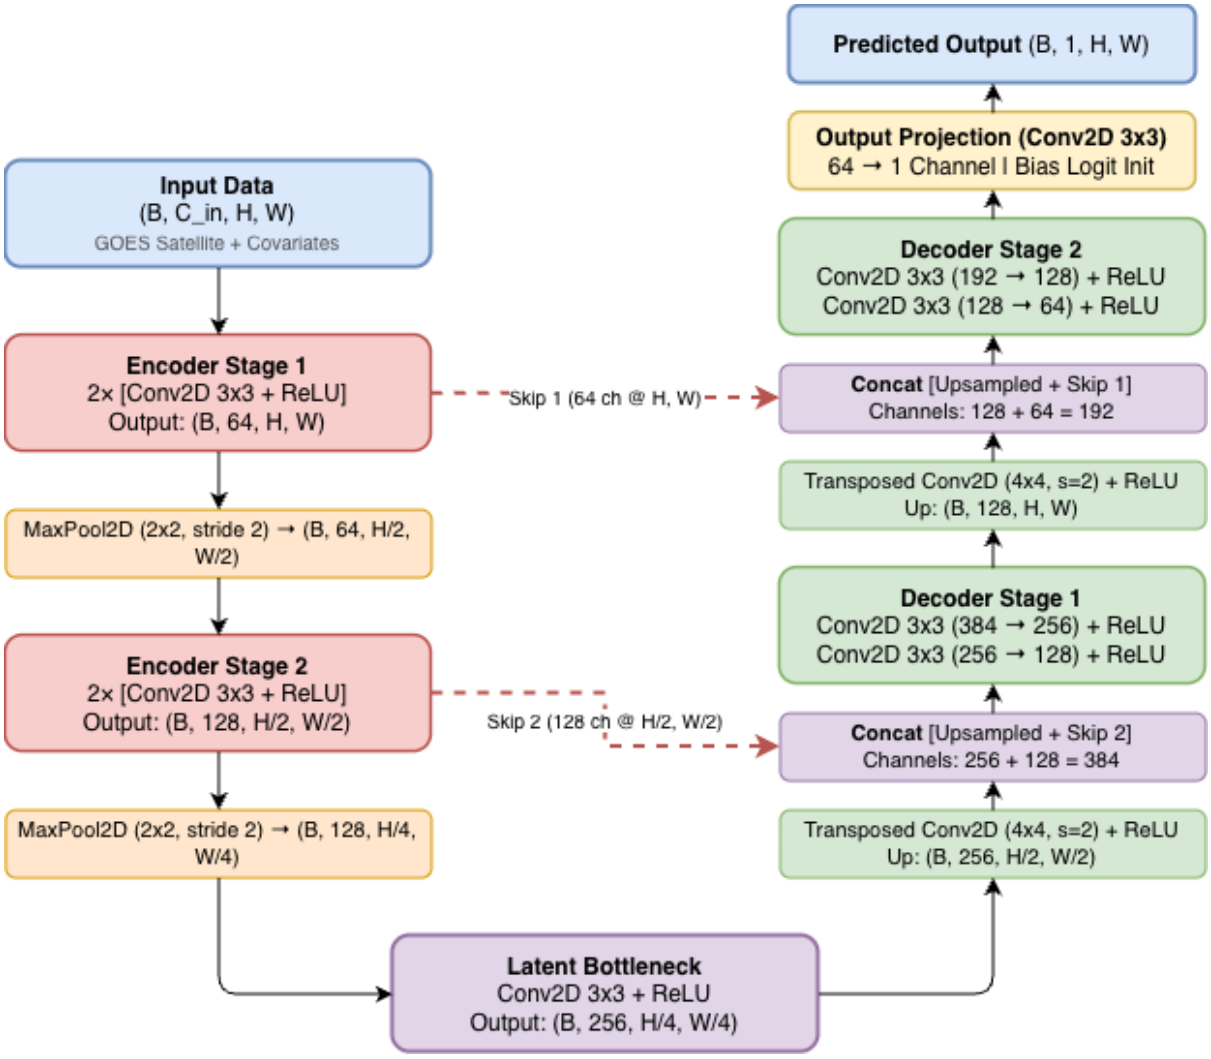


Figure 2: Architecture overview of the autoencoder baseline based on Badhan et al., 2024.

This baseline is trained under identical conditions to HAUNet (loss, optimiser, splits, covariates), isolating whether any performance gap is attributable to HAUNet's attention mechanisms rather than a simpler convolutional model.

4.3 Loss Functions

4.3.1 Focal Dice Loss

Active fire pixels represent a small fraction of each patch, resulting in extreme class imbalance. To address this, Focal Dice Loss was used (P. Wang and Chung, 2018). While Focal loss forces the model to focus on hard-to-classify edge pixels, the Dice component maximises spatial region overlap. It combines Focal Loss (which down-weights background pixels via $(1-p_t)^\gamma$) with Dice Loss:

$$\mathcal{L}_{\text{FD}} = w_f \cdot \mathcal{L}_{\text{focal}}(\alpha, \gamma) + w_d \cdot \mathcal{L}_{\text{dice}} \tag{3}$$

where w_f and w_d are weights of the losses, $\alpha \in [0, 1]$ balances positive and negative pixel contributions, and $\gamma \geq 0$ controls the rate at which pixels are down-weighted. Full formulations are given in A.1.

The hyperparameters used are: $\alpha = 0.85, \gamma = 1.5, w_f = 1.0, w_d = 1.0$, which was assigned for final model training after tuning with subset of training samples $(1,000)$. A higher α was assigned to penalise false negatives more heavily given the extreme class imbalance.

4.4 Training Setup

The dataset is split into training/validation/testing (80/10/10). This split uses stratified random sampling across four quantile bins based on the active pixel count for balanced representation. Weighted random sampling is applied during training to address class imbalance, where sample weights w_i for each patch i were assigned based on the count of positive fire pixels n_i:

$$w_i = \max(n_i, 1) \tag{4}$$

Note: this means training gradients are dominated by larger fires, and metrics are weighted towards performance on larger fires (McDermott et al., 2025). For Artio's use case this is not necessarily a limitation, as larger fires typically represent the greater share of insured exposure. However reported metrics should not be read as uniformly representative of performance on small fires.

Inputs are scaled prior to model ingestion. The GOES and DEM channels undergo Z-score standardisation using mean and standard deviation precomputed exclusively from training split. DEM covariates were converted to km. The landcover data is also one-hot encoded into 10 channels, each representing a landcover class.

For the targets, the continuous VIIRS FRP values are binarised, and Landsat targets are already binarised. Depending on the experimental configuration, these are transformed using spatial dilation.

Spatial dilation was applied using max-pooling kernel (3×3 for VIIRS and 7×7 for Landsat) to tackle class imbalance, giving denser signal for the model to learn (P. Wang and Chung, 2018). Landsat targets are much sparser than the VIIRS, so larger dilation kernel was used to provide a denser training signal.

The Landsat branch has limited number of samples $(1,202$ total), and given the extreme class imbalance in fire pixels, training samples were augmented via independent horizontal and vertical flips (each applied with probability 0.5), applied to the input and target. This increases the effective diversity of the training set. No augmentation was applied to validation or test splits.

All models were trained for 20 epochs using the Adam optimiser with initial learning rate $\eta = 1 \times 10^{-4}$, decaying to a minimum learning rate of 1×10^{-6}, and Weight decay 5×10^{-4} for HAUNet and 1×10^{-4} for the Autoencoder. Batch size of 8 were used to accommodate high-resolution multi-channel tensors within GPU VRAM constraints without downsampling spatial patch dimensions.

4.5 Computational Environment

Model training were conducted on a GCP Vertex AI Workbench instance with NVIDIA L4 GPU, 23.7 GB VRAM, 32 GB RAM, 8 vCPUs. Data generation and preprocessing were run as separate CPU-only jobs.

4.6 Experimental Design

HAUNet was assessed under the following ablations, in addition to the direct-vs-cascaded pipeline comparison:

1. with and without the Covariate Pyramid Encoder.
2. with and without DEM as an input covariate.
3. with and without land cover as an input covariate.
4. with and without target dilation (to assess the effect of target representation on model performance).

In the cascaded pipeline, Stage 1 (GOES $\rightarrow$ VIIRS) was trained on GOES/VIIRS pairs to predict VIIRS-resolution (375m) fire mask from GOES input. Stage 2 (VIIRS $\rightarrow$ Landsat) is intended to further downscale this to Landsat resolution (30m). However, VIIRS and Landsat overpasses times do not overlap, so matching VIIRS/Landsat pairs are not available in the same way GOES/VIIRS pairs are. To address this mismatch, Stage 2 is trained using the trained Stage 1 model's predictions on GOES/Landsat dataset as inputs. For each GOES/Landsat pair, the Stage 1 model produces a predicted VIIRS-resolution fire mask from the temporally matched GOES input, and this prediction is paired with the corresponding Landsat ground truth to train Stage 2. This ensures both stages of the cascaded pipeline operate on inputs available at inference time (i.e. Stage 2 never has access to real VIIRS observations at test time either), avoiding train/inference mismatch that would occur if Stage 2 were trained on ground-truth VIIRS data unavailable in a genuinely cascaded operational setting.

Dilation was investigated for potentially mitigating the severe class imbalance to help the model to learn fire regions more effectively. However, dilation can blur the boundaries of fire regions, potentially leading to less precise predictions. The effectiveness of this approach was evaluated through experiments comparing models trained with and without dilation. Dilation is performed by applying max pooling to the target masks. For this experiment, model training and validation is conducted on dilated targets, but evaluation is performed on the original, undilated targets.

Each experiment was defined by YAML files specifying model hyperparameters, data paths, loss weights, and training settings for reproducibility.

4.7 Evaluation Metrics

For Stage 1 cascaded pipeline, to align the model prediction dimensions (256×256) with the target (20×20), an adaptive spatial max-pooling was applied to the output logits before calculating loss and metrics. All reported quantitative metrics reflect performance evaluated on the coarse 20×20 target.

- **AUPRC (Area Under Precision-Recall Curve):** Measures the performance of the model across different thresholds, making it suitable for imbalanced datasets (Davis and Goadrich, 2006):

$$\text{AUPRC} = \sum_k (R_k - R_{k-1}) P_k \tag{5}$$

Table 3: Summary of experimental configurations evaluated.

Exp.	**Model**	**Pipeline**	**Exp Config.**	**Purpose**
E1	HAUNet	Direct	Full (with pyramid encoder, no dilation)	Primary configuration
E1a	HAUNet	Direct	No pyramid encoder	Ablation: pyramid branch contribution
E1b	HAUNet	Direct	No DEM	Ablation: DEM contribution
E1c	HAUNet	Direct	No land cover	Ablation: land cover contribution
E1d	HAUNet	Direct	Full, dilated target (k=7)	Ablation: target representation (dilated vs. default)
E2.1a	HAUNet	Cascaded (S1)	Full, no dilation	Direct vs. cascaded
E2.1b	HAUNet	Cascaded (S1)	Full, dilated targets (k=3)	Ablation: target representation
E2.2a	HAUNet	Cascaded (S2)	Best config. from E1, E2.1a input	Direct vs. cascaded
E2.2b	HAUNet	Cascaded (S2)	Best config. from E1, E2.1b input	Ablation: target representation
E3	Autoencoder	Direct	Full	Baseline comparison
E3d	Autoencoder	Direct	Full, dilated target	Baseline comparison, matched to E1d
E4	Autoencoder	Cascaded	Full, dilated and undilated target	Baseline, cascaded pipeline

where P_k and R_k represent precision and recall evaluated at the k-th threshold operating point. This is selected as primary metric for evaluation as it is threshold-independent and less sensitive to extreme class imbalance. It captures the trade-off between precision and recall across all thresholds.

- **IOU (Intersection over Union):** Measures the overlap between predicted and ground truth. Given a binary prediction mask $\hat{Y} \in \{0,1\}^{H\times W}$ and binary target mask $Y \in \{0,1\}^{H\times W}$:

$$\mathrm{IoU}(\hat{Y}, Y) = \frac{|\hat{Y} \cap Y|}{|\hat{Y} \cup Y|} = \frac{\sum_{i,j} \hat{Y}_{ij} Y_{ij}}{\sum_{i,j} \left[\hat{Y}_{ij} + Y_{ij} - \hat{Y}_{ij} Y_{ij}\right]} \tag{6}$$

- **Boundary F1:** Restricts precision and recall to fire-perimeter only, obtained by morphological erosion of the predicted and ground truth masks. This isolates edge/shape accuracy from interior-region, which pixel-level F1 can mask:

$$\mathrm{F1}_b = 2 \cdot \frac{\mathrm{Prec}_b \cdot \mathrm{Rec}_b}{\mathrm{Prec}_b + \mathrm{Rec}_b + \epsilon} \tag{7}$$

  where Prec_b and Rec_b are precision and recall computed only over boundary pixels of the predicted and ground truth masks.

- **Precision and Recall:** Evaluates the model's ability to correctly identify fire pixels (precision) and capture all true fire pixels (recall) (Saito and Rehmsmeier, 2015):

$$\text{Precision} = \frac{TP}{TP + FP + \epsilon}, \quad \text{Recall} = \frac{TP}{TP + FN + \epsilon} \tag{8}$$

  Where TP, FP, and FN represent true positive, false positive, and false negative counts across the spatial grid, and $\epsilon = 10^{-8}$ prevents division by zero.

For Boundary F1, IOU, Precision, and Recall, the evaluation is performed on post-processed outputs using Otsu's thresholding to binarise the predicted downscaled masks (Zhang and Hu, 2008). This calculates the variance between the foreground and background classes to determine an optimal threshold t^* that maximises the separation between them. Given class probabilities $\omega_0(t), \omega_1(t)$ and means $\mu_0(t), \mu_1(t)$ for pixels above/below threshold t, the variance is:

$$\sigma_B^2(t) = \omega_0(t)\,\omega_1(t)\,\left(\mu_0(t) - \mu_1(t)\right)^2, \qquad t^* = \arg\max_t\ \sigma_B^2(t) \tag{9}$$

Otsu threshold is calculated and applied per patch. Otsu is used over a fixed threshold (e.g. 0.5) because the extreme class imbalance means that the model may predict low probabilities for the positive class, and a fixed threshold could lead to poor separation.

To find uncertainty of metrics given the limited test-set size, 95% confidence intervals (CI) were computed on AUPRC via percentile bootstrap resampling at the patch level (1,000 resamples with replacement). To assess whether differences between configurations were statistically meaningful, paired bootstrap tests were performed for selected comparisons: same resampled patches were used to compute AUPRC for both configurations in each resample, giving distribution of paired differences from which a 95% CI on the difference, and find the proportion of resamples favouring one configuration.

Due to the extreme rarity of the positive class, absolute metric values are difficult to interpret without a reference point. The following baselines were used:

1. an all-background predictor, assigning zero probability to every pixel.

2. a constant-prevalence predictor, assigning every pixel the measured test-set fire-pixel prevalence.

2. is the theoretically expected value of AUPRC under a classifier with no discriminative ability. Model performance is reported alongside both baselines.

While target dilation were used for the training target to mitigate class imbalance during optimisation, all evaluation metrics were computed against the original, undilated binary target masks. This ensures reported performance reflects the model's ability to predict at native target resolution.

The best model checkpoint during fine tuning and selection for final evaluation and was selected based on the highest validation AUPRC. AUPRC evaluates across range of classification thresholds, making it threshold-independent and less sensitive to the extreme class imbalance.

5 Results

The evaluation results for the direct and cascaded pipelines are summarised in Table 4. All metrics in this study are aggregated (per-sample, then averaged across the test set). Example predictions and corresponding input plots for the best Autoencoer and HAUNet configuration are shown in A.3 and A.2. Figure 3 shows the Precision-Recall (PR) curves across representative models compared to the random baseline. Table 5 summarises the best-performing configuration within each architecture/pipeline group.

Table 4: Model performance evaluation across feature ablations vs. fixed baselines. Precision, Recall, Boundary F1 are computed on Otsu-thresholded predictions. Bold indicates the best value for that metric among the shown ablations of the same architecture and target

Exp.	**AUPRC**	**IoU**	**Precision**	**Recall**	**Boundary F1**
Direct (GOES->Landsat) - HAUNet					
E1: Full, No Target Dilation	0.0103	0.0021	0.0021	0.3403	0.0030
E1a: No Cov. Pyramid Enc.	0.0098	0.0043	0.0043	0.4605	0.0039
E1b: No DEM	0.0050	0.0020	0.0020	0.3100	0.0033
E1c: No Land Cover	0.0054	0.0026	0.0026	**0.5051**	0.0020
E1d: Dilated Target	**0.0188**	**0.0044**	**0.0044**	0.3490	**0.0041**
Cascaded (GOES->VIIRS->Landsat) - HAUNet					
E2.1a: Stage 1 (no dilation)	**0.1458**	**0.0651**	**0.0721**	0.3582	**0.0956**
E2.1b: Stage 1 (dilated)	0.1408	0.0505	0.0516	**0.6145**	0.0563
E2.2a: Stage 2 (E2.1a input, dilated)	**0.0142**	0.0072	0.0074	0.2065	0.0071
E2.2b: Stage 2 (E2.1b input dilated)	0.0134	**0.0089**	**0.0095**	**0.2435**	**0.0104**
Direct (GOES->Landsat) - Autoencoder					
E3: Full, No Target Dilation	**0.0120**	**0.0064**	**0.0077**	0.2592	**0.0050**
E3d: Dilated Target	0.0098	0.0057	0.0058	**0.5908**	0.0045
Cascaded pipeline - Autoencoder					
E4.1: Stage 1 (no dilation)	0.0944	0.0327	0.0367	0.3966	0.0500
E4.2a: Stage 2 (no dilation)	0.0106	**0.0065**	**0.0080**	0.2761	0.0037
E4.2b: Stage 2 (dilated targets)	**0.0106**	0.0056	0.0058	**0.4565**	**0.0040**
Baseline: All-Zero	0.0006	0.0000	0.0000	0.0000	N/A
Baseline: Prevalence-Rate	0.0006	0.0000	0.0000	0.0000	N/A

6 Discussions

6.1 Overall Performance

Across configurations, absolute performance is generally low (most < 0.02 AUPRC and IOU). But all exceed the All-Zero and Prevalence-Rate baselines, meaning the models are learning the fire signal rather than predicting background (relative to prevalence, best direct configuration (E1d) achieves an AUPRC of 0.0188, roughly 30$\times$ this baseline, the best cascaded configu-

Table 5: Best-performing configuration (by AUPRC) within each architecture/pipeline group, for direct cross-group comparison. Confidence intervals with with 95% bootstrap are shown in parentheses.

Group	AUPRC	IoU	Precision	Recall	Boundary F1
Direct - HAUNet (E1d: Dilated Target)	**0.0188 [0.0084, 0.0336]**	0.0044	0.0044	0.3490	0.0041
Cascaded - HAUNet (E2.2a: Stage 2 (E2.1a input, dilated))	0.0142 [0.0071, 0.0237]	**0.0072**	0.0074	0.2065	**0.0071**
Direct - Autoencoder (E3: Full, No Target Dilation)	0.0120 [0.0044, 0.0243]	0.0064	**0.0077**	0.2592	0.0050
Cascaded - Autoencoder (E4.2b: Stage 2 (dilated targets))	0.0106 [0.0044, 0.0184]	0.0056	0.0058	**0.4565**	0.0040

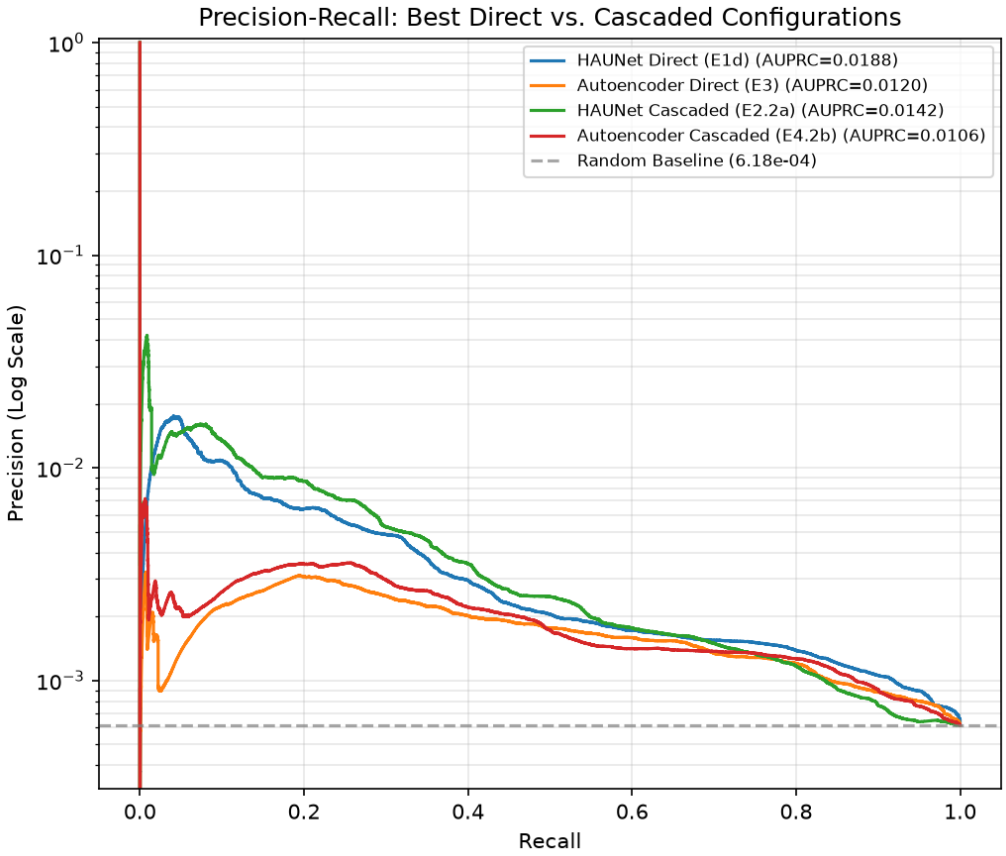


Figure 3: Precision-Recall curves for best Direct and Cascaded configurations (E1d, E2.2a, E3, E4.2b) against the random baseline (Prevalence-Rate = 6.18×10^{-4}). Curves evaluate raw probability outputs across all decision thresholds.

ration (E2.2a) roughly 23×, and the Autoencoders (E3, E4.2b) roughly 20× and 17× times respectively).

Direct comparison with Badhan et al., 2024 is not meaningful given differences in target resolution (2 km → 375 m vs. this study's 2 km → 30 m), prediction target (continuous BT vs. binary fire), and evaluation method. However, their IOU values (0.03-0.24) achieved with a single smaller resolution jump, less severe class imbalance (they use 20% fire-coverage threshold, and smaller image 128×128), and more target dataset (their $5,869$ samples vs. $1,202$ Landsat target samples in this study) still results in low overlap. Suggesting the extreme resolution gap and class imbalance in this study make the low absolute performance unsurprising.

95% CIs were computed for AUPRC via patch-level bootstrap resampling (Table 5). These are wide relative to the differences between configurations, reflecting the limited test-set size.

6.2 Over-Prediction, Thresholding, and Practical Implications

Across most configuration, Recall exceeds Precision (e.g. E1c: Recall 0.5051 vs. Precision 0.0026; E1d: Recall 0.3487 vs. Precision 0.0044). This indicates the models are over-predicting fire pixels. However, in insurance, over-predicting may be preferable to under-predicting, as false negatives could have more severe consequences than false positives. This pattern is also consistent with McCarthy et al., 2021, who report high recall with collapsing precision in their GOES-to-VIIRS U-Net, due to hot gases producing false thermal signal around the true fire area. It also raises whether Otsu is well-suited to thresholding predictions, as it assumes a roughly bimodal intensity distribution. A fixed threshold or one derived from validation PR curve may be worth comparing in future work. The Prevalence-Rate baseline, has no bimodal structure and Otsu collapses to the same all-negative mask as All-Zero baseline, hence identical.

To quantify over-prediction, area inflation ratio was computed (predicted fire area over true fire area). Median inflation ranging from 39× (E2.2a) to 457× (E4.2b) across the four configurations, amplified by extreme rarity of the positive class (Landsat fire pixels comprise only 0.06% of a patch on average). Note this is a pixel-count ratio than spatial-overlap, it should be read alongside the low IoU values, which already indicate limited spatial localisation.

For Artio's carbon investment risk assessment use case, under-estimating fire is generally more costly than overestimating: the current models are better suited to coarse regional estimates than precise predictions, and any operational use of predicted area would need calibration against this inflation factor. Notably, E2.2a shows lower inflation (39×) than the other configurations, suggesting the cascaded HAUNet may produce more spatially concentrated predictions. Though given the overlapping CIs, this should be treated as observation rather than a confirmed architectural advantage.

6.3 Cascaded vs. Direct Pipeline

For HAUNet, Stage 2 (E2.2a: IoU 0.0072, Precision 0.0074, AUPRC 0.0142 Boundary F1 0.0071, E2.2b: IOU 0.0089, Precision 0.0095, Boundary F1 0.0104) outperforms best direct configuration (E1d: IoU 0.0044, Precision 0.0044, AUPRC 0.0188, Boundary F1 0.0041) on all metrics except AUPRC and Recall. This suggests the intermediate stage may help the model produce spatially better localised predictions (reflected in overlap-based metrics) but not necessarily improve confidence ranking. Stage 1 metrics are much higher than either Stage 2 or direct-pipeline (e.g. E2.1a AUPRC 0.1458), but this is due to Stage 1 metric being evaluated at VIIRS target and resolution (20×20). So metrics are not directly comparable.

Paired bootstrap test on AUPRC (E2.2a vs. E1d, matched by patch, $1,000$ resamples) gives mean difference of -0.0042 with 95% CI $[-0.0192, 0.0086]$, and E2.2a exceeds E1d in only 30.2% of resamples. While not significant, it is consistent with the direct configuration retaining an AUPRC advantage, even as the cascaded configuration improves overlap-based metrics. Reinforcing that these two metric families capture different aspects of prediction quality.

For the Autoencoder, there is similar pattern. The direct configuration achieves highest AUPRC (E3: 0.0120), while cascaded Stage 2 achieve only slightly better IoU and Precision (E4.2a: IoU 0.0065, Precision 0.0080 vs. E3: IoU 0.0064, Precision 0.0077) and higher Recall (E4.2b: 0.4565 vs. E3: 0.2592). Errors from Stage 1 prediction propagate into Stage 2 training inputs, introducing noise that lowers confidence quality (AUPRC).

A paired bootstrap test on AUPRC (E4.2b vs. E3, matched by patch) gives a mean difference of -0.0013 with 95% CI $[-0.0082, 0.0052]$, and E4.2b exceeds E3 in only 32.3% of resamples, indicating a weaker and less consistent AUPRC gap for the Autoencoder than observed for HAUNet, though neither reaches statistical significance at this sample size.

The divergence between AUPRC and the overlap-based metrics (IoU, Precision, Boundary F1) are seen across comparisons. Because AUPRC is computed over continuous predictions and reflects how well the model ranks pixels by confidence across all thresholds. IoU, Precision, and Boundary F1, are computed only after Otsu thresholding. They depend on the binary mask produced only at one particular threshold. A model can improve performance at specific Otsu threshold (e.g. producing a more spatially coherent, better-localised binary region) without improving its overall confidence across all thresholds. This is a possible explanation for why the cascaded stage, and dilated inputs, improve overlap-based metrics for HAUNet while not improving AUPRC.

6.4 HAUNet vs. Autoencoder Architecture Comparison

Comparing best Stage 2 configurations, HAUNet achieves higher AUPRC point estimates than the Autoencoder in both pipelines:

- **Cascaded Pipeline:** HAUNet leads autoencoder on AUPRC (E2.2a: 0.0142 vs. E4.2a / E4.2b: 0.0106), IoU (E2.2b: 0.0089 vs. E4.2a: 0.0065), Precision (0.0095 vs. 0.008) and Boundary F1 (0.0104 vs. 0.004).
- **Direct Pipeline:** HAUNet leads Autoencoder on AUPRC (E1d: 0.0188 vs. E3: 0.0120), Recall (0.3490 vs. 0.2592). While direct Autoencoder achieves higher IoU (0.0064 vs. 0.0044) and Precision (0.0077 vs. 0.0044).

Paired bootstrap tests on AUPRC support this pattern: HAUNet exceeds the Autoencoder in 84.2% of resamples in the direct pipeline (E1d vs. E3, mean diff +0.0069, 95% CI $[-0.0057, 0.0222]$) and 87.1% in the cascaded pipeline (E2.2a vs. E4.2b, mean diff +0.0036, 95% CI $[-0.0028, 0.0110]$). Neither interval excludes zero at the conventional 95% threshold given the limited test-set size, but the consistency of this advantage across both pipelines and across roughly six in seven resamples suggests architectural difference favouring HAUNet.

Autoencoder predictions produce broader, low-confidence regions, likely due to their convolutional nature. After Otsu thresholding, these wider predictions cover larger ground areas, yielding higher point precision/IoU at the cost of spatial localisation. HAUNet produces sharper predictions with higher overall confidence ranking (PR curves in 3). Furthermore, HAUNet Stage 1 provides higher-quality inputs to Stage 2 (E2.1a AUPRC 0.1458 vs. E4.1 AUPRC 0.0944) meaning less severe error propagation in the cascaded pipeline. Suggesting the results differ by Stage 2 architecture's own suitability for the task, and differences in the quality of Stage 1 predictions. Investigating these effects require cross-architecture ablation (e.g. evaluating an Autoencoder Stage 2 trained on HAUNet's Stage 1 predictions, vice versa), which is left for future work.

6.5 Effect of Target Dilation

Dilating the target improves AUPRC and IoU for HAUNet's direct configuration (E1d vs. E1: AUPRC 0.0188 vs 0.0103, IoU 0.0044 vs 0.0021) but not for the Autoencoder (E3d vs. E3:

AUPRC 0.0098 vs 0.0120, IoU 0.0057 vs 0.0064, both lower with dilation). This may be because dilation allows for smoothing/tolerance for models like HAUNet that uses attention mechanisms to capture fine-grained information, while Autoencoder may benefit less from this because their convolutional nature smoothes outputs. This is consistent with the higher recall.

Note that Stage 1 dilation ablation (E2.1a vs. E2.1b) uses smaller kernel matched to the coarser VIIRS grid ($k = 3$ vs. $k = 7$) and shows AUPRC decrease with dilation (0.1458 vs. 0.1408). This is likely explained by the relative size of the dilation kernel: a 3×3 kernel on the 20×20 VIIRS grid affects approximately 15% of the grid dimension, compared to a 7×7 kernel on the 256×256 Landsat grid, which affects only around 2.7%. At VIIRS resolution, dilation of this relative magnitude may distort the shape of an already coarse, fire signal rather than providing the tolerance it offers at Landsat resolution. This suggests the appropriate dilation kernel size should be chosen relative to target grid resolution, and that the direct-pipeline dilation effect is not directly comparable to Stage 1.

6.6 Feature Ablations

Removing DEM (E1b) produces the weakest performance (AUPRC 0.0050, IoU 0.0020 vs. E1: 0.0103, 0.0021), and removing land cover (E1c) causes AUPRC drop (0.0054) alongside the group's highest Recall (0.5051) but lowest Boundary F1 (0.0020). Together suggesting both covariates contribute to spatial precision.

Removing the covariate pyramid encoder (E1a) unexpectedly improves overlap-based metrics relative to the full model (IoU: 0.0043 vs. 0.0021; Precision: 0.0043 vs. 0.0021; Boundary F1: 0.0039 vs. 0.0030), while AUPRC is only slightly lower (0.0098 vs. 0.0103). One possible explanation is that without full-resolution covariate detail to sharpen its predictions, the model may default to coarser, more spatially diffuse output, similar in effect to the Autoencoder's convolutional smoothing discussed above. Broader predictions can improve overlap metrics, without necessarily improving the model's overall confidence ranking (reflected in the near-unchanged AUPRC). This may suggest the pyramid encoder's contribution is in sharpening prediction confidence and preserving fine spatial detail rather than in overall discriminative ability. This is worth confirming in future work.

7 Conclusions

This project investigated DL pipeline for downscaling GOES imagery into Landsat-resolution wildfire predictions, comparing direct and cascaded mapping across HAUNet and Autoencoder baseline, in the context of Artio's carbon-investment risk assessment use case.

Absolute performance remained modest (AUPRC and IoU < 0.02), reflecting the extreme resolution gap (2 km to 30 m), severe class imbalance, and limited training data ($1,202$ Landsat patches). All configurations exceeded all-zero and prevalence-rate baseline (AUPRC up to $30\times$ the random-classifier baseline) confirming genuine signal extraction rather than deployment-ready accuracy.

No pipeline design showed uniform advantage, and architecture comparisons were pipeline-dependent: in the direct pipeline the Autoencoder achieved better fixed-threshold overlap metrics while HAUNet led on AUPRC and Recall, the cascaded pipeline HAUNet led on every metric. Paired bootstrap tests support consistent AUPRC advantage for HAUNet across both pipelines, though this was not statistically significant at this sample size. This divergence, alongside a large area inflation ratio (up to $457\times$ at Otsu threshold), suggests the current models are better suited to coarse regional risk flagging than precise active fire prediction for Artio's use case. Which has operational relevance, since missed fire is costlier than over-prediction.

Future work should prioritise a cross-architecture ablation isolating Stage 2 performance from Stage 1 input quality, and integration of temporal or climate covariates to move toward

operational active fire and intensity prediction.

A Appendix

A.1 Focal Dice Loss Formulation

Focal Loss (Lin et al., 2018) applies $(1 - p_t)^\gamma$ to focus training on sparse positives:

$$\mathcal{L}_{\text{focal}} = -\alpha_t (1 - p_t)^\gamma \log(p_t) \tag{10}$$

where α_t is the balancing factor, p_t is the predicted probability for the ground-truth class, and $p_{i,j} = P(y_{i,j} = 1)$ is the estimated probability for the positive class. For a ground-truth label $y_{i,j} \in \{0, 1\}$:

$$p_t = \begin{cases} p_{i,j} & \text{if } y_{i,j} = 1 \\ 1 - p_{i,j} & \text{if } y_{i,j} = 0 \end{cases}, \quad \alpha_t = \begin{cases} \alpha & \text{if } y_{i,j} = 1 \\ 1 - \alpha & \text{if } y_{i,j} = 0 \end{cases} \tag{11}$$

$\alpha \in [0, 1]$ balances positive and negative pixel contributions, and $\gamma \geq 0$ controls the rate at which pixels are down-weighted.

Dice Loss (Milletari et al., 2016) tries to maximise the spatial overlap:

$$\mathcal{L}_{\text{dice}} = 1 - \frac{2 \sum_{i,j} p_{i,j} y_{i,j} + \epsilon}{\sum_{i,j} p_{i,j} + \sum_{i,j} y_{i,j} + \epsilon} \tag{12}$$

where ϵ is a smoothing term preventing division by zero.

A.2 Input GOES and Covariate channels

Visual example of the input GOES and covariate channels (inputs used for examples shown in A.3 are shown in Figure 4.

A.3 Prediction Plots

Figures 5 - 8 show example predictions for best direct and cascaded HAUNet and Autoencoder models.

References

Badhan, M., Shamsaei, K., Ebrahimian, H., Bebis, G., Lareau, N. P., & Rowell, E. (2024). Deep Learning Approach to Improve Spatial Resolution of GOES-17 Wildfire Boundaries Using VIIRS Satellite Data. *Remote Sensing*, *16*(4), 715. https://doi.org/10.3390/rs16040715

Davis, J., & Goadrich, M. (2006). The relationship between Precision-Recall and ROC curves. *Proceedings of the 23rd International Conference on Machine Learning*, 233–240. https://doi.org/10.1145/1143844.1143874

de Almeida Pereira, G. H., Fusioka, A. M., Nassu, B. T., & Minetto, R. (2021). Active fire detection in Landsat-8 imagery: A large-scale dataset and a deep-learning study. *ISPRS Journal of Photogrammetry and Remote Sensing*, *178*, 171–186. https://doi.org/10.1016/j.isprsjprs.2021.06.002

Dwyer, O., & Nam, H. W. (2026, March 31). *How wildfires and storms drove insurance losses in 2025 – in three charts*. Carbon Brief. Retrieved June 10, 2026, from https://www.carbonbrief.org/how-wildfires-and-storms-drove-insurance-losses-in-2025-in-three-charts/

Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D., & Moore, R. (2017). Google Earth Engine: Planetary-scale geospatial analysis for everyone. *Remote Sensing of Environment*, *202*, 18–27. https://doi.org/10.1016/j.rse.2017.06.031

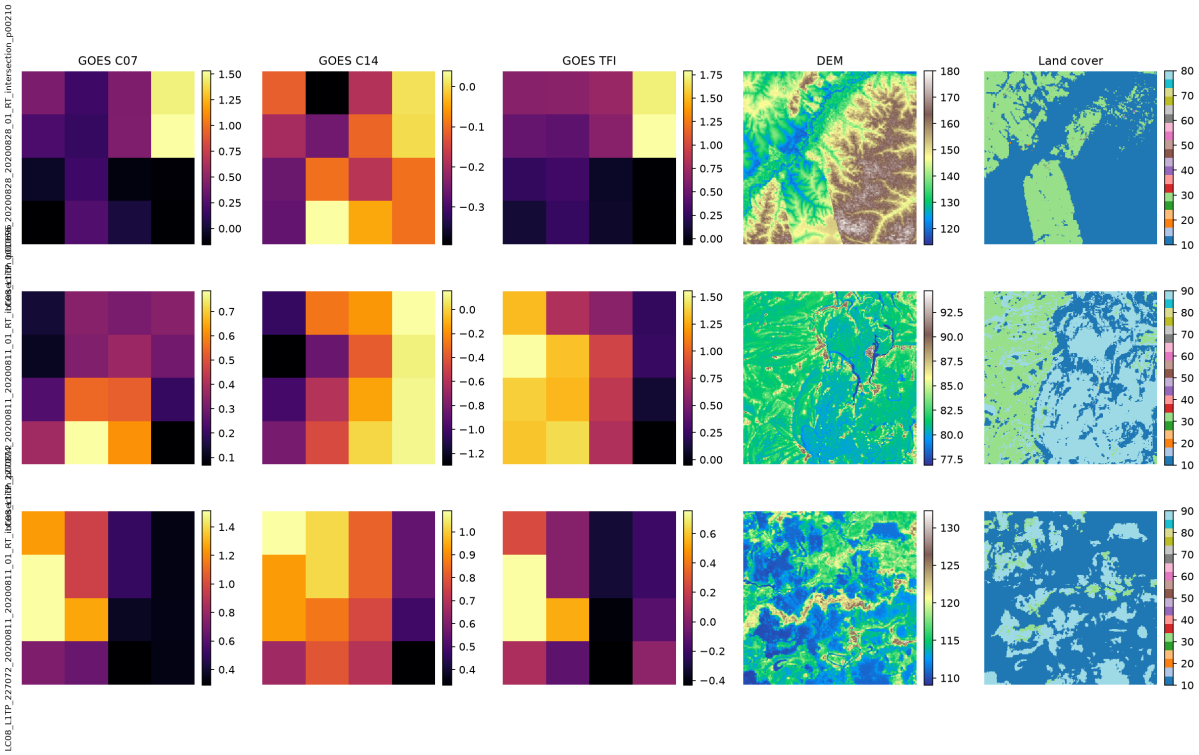


Figure 4: GOES and covariates inputs for same samples as figures in A.3. Left to right: GOES C07, C14, TFI, DEM, Landcover.

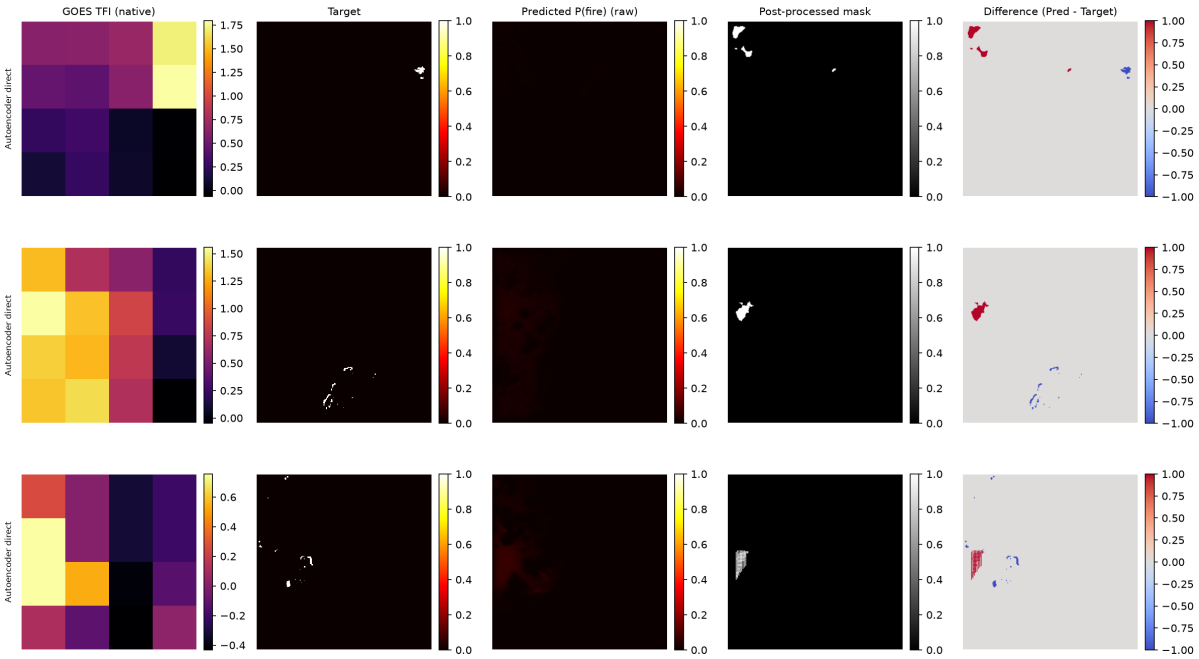


Figure 5: Autoencoder direct (E3) prediction examples. From left to right: GOES TFI, undilated target, dilated target, raw predictions, postprocessed predictions, difference.

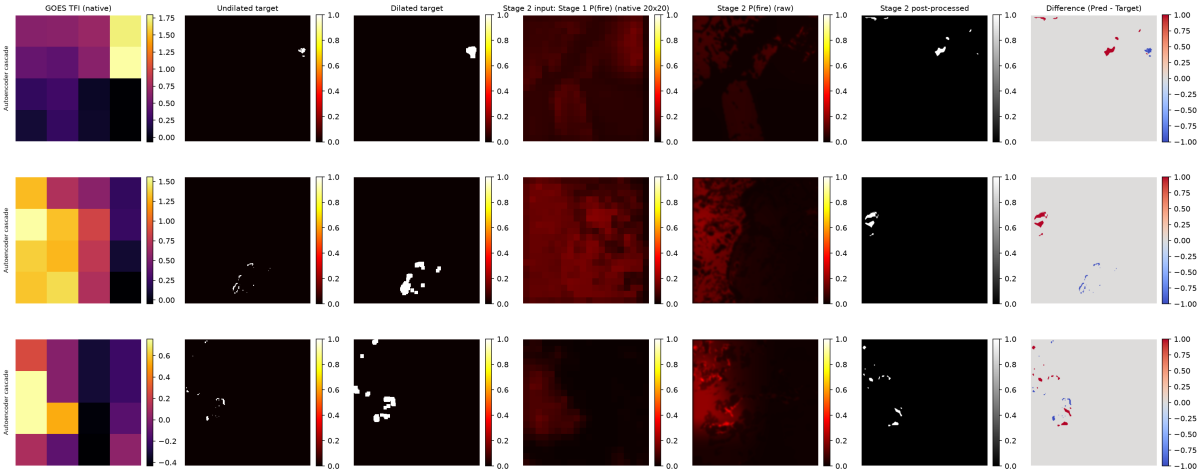


Figure 6: Autoencoder cascaded (E4.2b) prediction examples. From left to right: GOES TFI, undilated target, dilated target, raw Stage 1 predictions, raw Stage 2 predictions, postprocessed predictions, difference.

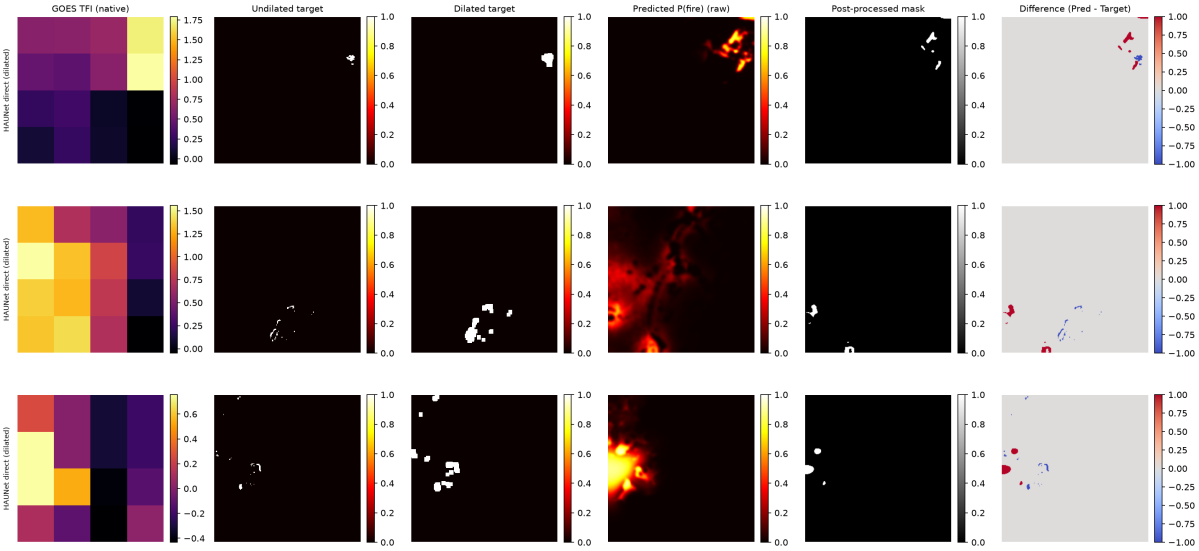


Figure 7: HAUNet direct (E1d) prediction examples. From left to right: GOES TFI, undilated target, dilated target, raw predictions, postprocessed predictions, difference.

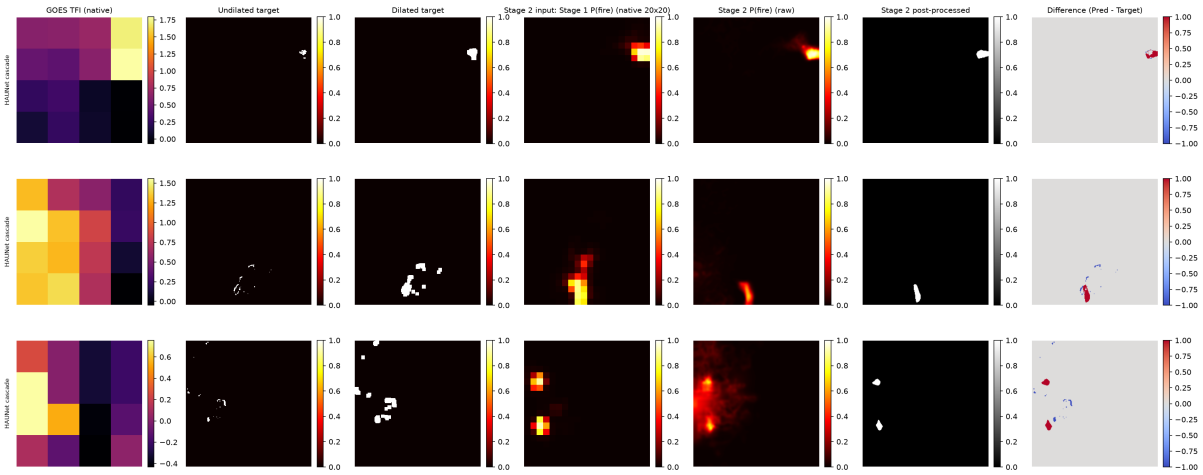


Figure 8: HAUNet cascaded (E2.2a) prediction examples. From left to right: GOES TFI, undilated target, dilated target, raw Stage 1 predictions, raw Stage 2 predictions, postprocessed predictions, difference.

Grassi, G., Monni, S., Federici, S., Achard, F., & Mollicone, D. (2008). Applying the conservativeness principle to REDD to deal with the uncertainties of the estimates. *Environmental Research Letters*, *3*(3), 035005. https://doi.org/10.1088/1748-9326/3/3/035005

Jones, M. W., Smith, A. J. P., Betts, R., Canadell, J. G., Prentice, I. C., & Quéré, C. L. (2020). Climate Change Increases the Risk of Wildfires.

Lambert, J., & Turner, G. (2024, November 14). *Investment Trends and Outcomes in the Global Carbon-Credit Market-2024 | MSCI*. Retrieved June 5, 2026, from https://www.msci.com/research-and-insights/paper/investment-trends-and-outcomes-in-the-global-carbon-credit-market-2024

Li, F., Zhang, X., Kondragunta, S., Schmidt, C. C., & Holmes, C. D. (2020). *A preliminary evaluation of GOES-16 active fire product using Landsat-8 and VIIRS active fire data, and ground-based prescribed fire records*. Retrieved August 21, 2026, from https://repository.library.noaa.gov

Li, X., Zhang, G., Zhu, S., & Xu, Y. (2025). Stepwise Downscaling of ERA5-Land Reanalysis Air Temperature: A Case Study in Nanjing, China. *Remote Sensing*, *17*(12), 2063. https://doi.org/10.3390/rs17122063

Lin, T.-Y., Goyal, P., Girshick, R., He, K., & Dollár, P. (2018, February 7). *Focal Loss for Dense Object Detection*. arXiv: 1708.02002 [cs.CV]. https://doi.org/10.48550/arXiv.1708.02002

Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., Lin, S., & Guo, B. (2021). Swin Transformer: Hierarchical Vision Transformer using Shifted Windows. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 9992–10002. https://doi.org/10.1109/ICCV48922.2021.00986

McCarthy, N. F., Tohidi, A., Aziz, Y., Dennie, M., Valero, M. M., & Hu, N. (2021). A Deep Learning Approach to Downscale Geostationary Satellite Imagery for Decision Support in High Impact Wildfires. *Forests*, *12*(3), 294. https://doi.org/10.3390/f12030294

McDermott, M. B. A., Zhang, H., Hansen, L. H., Angelotti, G., & Gallifant, J. (2025, January 13). *A Closer Look at AUROC and AUPRC under Class Imbalance*. arXiv: 2401.06091 [cs.LG]. https://doi.org/10.48550/arXiv.2401.06091

Milletari, F., Navab, N., & Ahmadi, S.-A. (2016, June 15). *V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation*. arXiv: 1606.04797 [cs.CV]. https://doi.org/10.48550/arXiv.1606.04797

NOAA. (2017, May 23). *GOES-16 Satellite Data Public Bucket*. https://console.cloud.google.com/storage/browser/gcp-public-data-goes-16

NOAA. (2018, November 14). *GOES-17 Satellite Data Public Bucket*. https://console.cloud.google.com/storage/browser/gcp-public-data-goes-17

Qi, Y., Lou, M., Liu, Y., Li, L., Yang, Z., & Nie, W. (2026). Advancing image super-resolution techniques in remote sensing: A comprehensive survey. *ISPRS Journal of Photogrammetry and Remote Sensing*, *231*, 68–100. https://doi.org/10.1016/j.isprsjprs.2025.10.024

Saito, T., & Rehmsmeier, M. (2015). The Precision-Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets. *PLOS ONE*, *10*(3), e0118432. https://doi.org/10.1371/journal.pone.0118432

Schmidt, C. C., Ssec, U.-M., Hoffman, J., Ssec, U.-M., Prins, E., Ssec, U.-M., Lindstrom, S., & Ssec, U.-M. (2013). GOES-R Advanced Baseline Imager (ABI) Algorithm Theoretical Basis Document For Fire / Hot Spot Characterization.

Sdraka, M., Papoutsis, I., Psomas, B., Vlachos, K., Ioannidis, K., Karantzalos, K., Gialampoukidis, I., & Vrochidis, S. (2022). Deep Learning for Downscaling Remote Sensing Images: Fusion and super-resolution. *IEEE Geoscience and Remote Sensing Magazine*, *10*(3), 202–255. https://doi.org/10.1109/MGRS.2022.3171836

Sullivan, A. L. (2017). Inside the Inferno: Fundamental Processes of Wildland Fire Behaviour. *Current Forestry Reports*, *3*(2), 150–171. https://doi.org/10.1007/s40725-017-0058-z

Wang, J., Chen, M., Karaev, N., Vedaldi, A., Rupprecht, C., & Novotny, D. (2025, March 14). *VGGT: Visual Geometry Grounded Transformer*. arXiv: 2503.11651 [cs.CV]. https://doi.org/10.48550/arXiv.2503.11651

Wang, J., Wang, B., Wang, X., Zhao, Y., & Long, T. (2023). Hybrid Attention-Based U-Shaped Network for Remote Sensing Image Super-Resolution. *IEEE Transactions on Geoscience and Remote Sensing*, *61*, 1–15. https://doi.org/10.1109/TGRS.2023.3283769

Wang, P., & Chung, A. C. S. (2018). Focal Dice Loss and Image Dilation for Brain Tumor Segmentation. In D. Stoyanov, Z. Taylor, G. Carneiro, T. Syeda-Mahmood, A. Martel, L. Maier-Hein, J. M. R. Tavares, A. Bradley, J. P. Papa, V. Belagiannis, J. C. Nascimento, Z. Lu, S. Conjeti, M. Moradi, H. Greenspan, & A. Madabhushi (Eds.), *Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support* (pp. 119–127). Springer International Publishing. https://doi.org/10.1007/978-3-030-00889-5_14

Wang, Y., Yang, W., Chen, X., Wang, Y., Guo, L., Chau, L.-P., Liu, Z., Qiao, Y., Kot, A. C., & Wen, B. (2023, November 23). *SinSR: Diffusion-Based Image Super-Resolution in a Single Step*. arXiv: 2311.14760 [cs.CV]. https://doi.org/10.48550/arXiv.2311.14760

Wang, Z., Chen, J., & Hoi, S. C. H. (2020, February 8). *Deep Learning for Image Super-resolution: A Survey*. arXiv: 1902.06068 [cs.CV]. https://doi.org/10.48550/arXiv.1902.06068

Zhang, J., & Hu, J. (2008). Image Segmentation Based on 2D Otsu Method with Histogram Analysis. *2008 International Conference on Computer Science and Software Engineering*, *6*, 105–108. https://doi.org/10.1109/CSSE.2008.206