Imperial College London
Department of Earth Science and Engineering
MSc Environmental Data Science and Machine Learning

Comparison of Direct and Cascaded High-Resolution Spatial Downscaling of Wildfire Data

Author:
Tsuyu Huang
tsuyu.huang25@imperial.ac.uk

Supervisors:
Jorge Veiras (jorge@artiocarbon.com)
Prof. Anna Korre (a.korre@imperial.ac.uk)

Code Repository:
https://github.com/ese-ada-lovelace-2025/irp-th25

Company:
Artio
www.artiocarbon.com

July 24, 2026

Abstract

With climate change, there's an increasing need for a higher-resolution wildfire data for better monitoring and risk assessment. GEO satellites have high-temporal resolution, but lacks spatial resolution for a detailed risk assessment. While LEO satellites provide better spatial resolution, but low temporal resolution to monitor rapid wildfire spread. Existing downscaling models have demonstrated their effectiveness in producing high-resolution remote sensing data. But many still fail to capture the fine details and physically plausible outputs. This project aims to achieve a high-resolution spatial downscaling of wildfire data by integrating additional physical covariates to the model to mitigate implausible outputs. The project explores direct (2km to 30m) and cascaded (2km to 375m to 30m) downscaling pipelines and their effectiveness in producing high-resolution wildfire data. The project compares diffusion (EDiffSR), attention U-Net (HAUNet), and transformer-based (ESTNet) downscaling models.

Contents

1 Introduction **3**

2 Current State of Remote Sensing Downscaling **3**

3 Data and Materials **4**

4 Methods **4**
4.1 HAUNet (Hybrid Attention U-Net) 4
4.1.1 Hybrid Attention Mechanisms 4
4.1.2 Full-Resolution Covariate Pyramid Branch 7
4.1.3 Multi-Scale Refinement and Feature Fusion 7

5 Results **7**

6 Discussions **7**

7 Conclusions **7**

1 Introduction

In recent years, due to rising popularity in Net-Zero initiatives, carbon investors are increasingly funding carbon reduction and removal projects, such as reforestation (Lambert and Turner, 2024). However, these carbon projects are at risk because wildfires can destroy and release the carbon back into the atmosphere. With climate change, wildfires are becoming increasingly severe and frequent (Jones et al., 2020). This causes economic losses for the insurance industry. For instance, The January 2025 Los Angeles County wildfires set an all-time record, causing $40 billion in insured losses alone from a single event (Dwyer and Nam, 2026). This project, conducted with Artio, aims to address this gap. By providing highly accurate risk assessment models, this research seeks to give better confidence and enable better risk assessment.

High-spatial and high-temporal resolution wildfire data is favourable in risk assessment. However, high-temporal resolution data obtained from geostationary satellites (GEO) have limited spatial resolution, which makes them difficult to use for accurate risk assessment. While low orbit satellites (LEO) offers high-spatial resolution, but low-temporal resolution.

To bridge this gap, existing methods uses Deep Learning (DL) models for downscaling due to their ability in recognising non-linear spatiotemporal patterns. However, the resolution gap between GEO and LEO is significant (2km - 30m). This increases the risk of producing physically implausible or unrealistic outputs, as a single low-resolution (LR) pixel can correspond to different high-resolution (HR) details (Z. Wang et al., 2020). In this project, to mitigate this limitation, the model integrates additional physical information, such as HR topography and landcover to help guide the downscaling process. To further improve downscaling, the project aims to investigate whether cascaded downscaling (2km to 375m to 30m) produces more accurate and better results compared to direct downscaling (2km to 30m). By breaking up the downscaling process, the model trains on a smaller resolution gap, aiming to mitigate physically implausible outputs. This project will compare the cascaded and direct pipeline performance of different downscaling architectures, including hybrid attention U-Net, a transformer-based model, and a diffusion-based model.

2 Current State of Remote Sensing Downscaling

Traditional downscaling methods often fail to capture the complex non-linear patterns and variability in remote sensing data. In recent years, DL methods have shown promising results that outpaces traditional approaches (Sdraka et al., 2022). In particular, Convolutional Neural Networks (CNNs) and U-Net architectures have been widely used for their ability to capture spatial features. Generative Adversarial Networks (GANs) have also been used for their ability to generate sharp images. Recently, Transformer and Diffusion-based approaches have been proposed for downscaling. Transformers use attention mechanisms to capture long-range dependencies (J. Wang et al., 2025), while Diffusion models shows good generative ability and fidelity (Y. Wang et al., 2023).

However, many of these approaches struggles preserving fine textures and geometric structures (Qi et al., 2026). This is problematic for wildfire data, where the spatial patterns and textures are crucial for risk assessment. To address these limitations, there is a need for domain-specific architectures that can better capture the characteristics of wildfire data.

Existing literature remains reliant on direct downscaling methods across large resolution-gaps. Some studies have explored cascaded approaches, such as AMCN-Geoi-DBN from Jing et al., 2022, and CLIT from Chen et al., 2023, but there are currently no research from the literature investigating cascaded downscaling in the context of wildfires. The limitations and benefits of cascaded downscaling is also largely unexplored. To address this gap, the project objectives are:

1. Develop a domain-specific, multi-stage framework for wildfire downsclaing.

2. Compare the performance of different downscaling architectures in terms of preserving fine textures and geometric structures.

3. Evaluate the trade-offs of cascaded and direct downscaling and its effectiveness in improving the accuracy of wildfire downscaling.

3 Data and Materials

Table 1 shows the data sources used for this project. The specifications for the VIIRS, GOES, and Landsat data are outlined in Table 2.

The data was processed so that all regions of interest (ROIs) are dynamically reprojected to their localised UTM zone based on centroid coordinates. Detections from the NASA VIIRS S-NPP archive were filtered for "nominal" and "high" confidence points. Detections are clustered into 7680×7680 m grid cells within a 15-minute temporal window of peak activity. VIIRS grid cell pixels record the peak Fire Radiative Power (FRP). Due to extreme class imbalance of the target fire datasets, stratified sampling was applied across temporal extent and fire pixel count bins during data splitting. The closest temporal GOES-16/17 scan within a $\pm$60-minute guard window was queried. In addition to the native 3.9 μm (Ch. 7) and 11.2 μm (Ch. 14) GOES brightness temperatures, a Thermal Fire Index (TFI) is computed as:

$$TFI = Ch.7 - Ch.14 \tag{1}$$

The dataset is split into training/validation/testing (80/10/10). This split utilises stratified random sampling across four quantile bins based on the active fire pixel count per patch ofr balanced representation of fire sizes.

Inputs are scaled prior to model ingestion. The GOES channels undergo Z-score standardisation using mean and standard deviation parameters precomputed exclusively from the training split.

For the targets, the continuous VIIRS FRP values are binarised into active fire masks. Depending on the experimental configuration, thes are subsequently transformed using either spatial dilation via a 3×3 max pooling operation or conversion into continuous Gaussian heatmaps.

4 Methods

4.1 HAUNet (Hybrid Attention U-Net)

Among CNN-based approaches, HAUNet (J. Wang et al., 2023) uses hybrid attention to capture both local fine detail and broader contextual structure across multiple scales. HAUNet was selected for its ability to capture local extreme hotspots in near-zero background, which standard U-Nets might smooth out. HAUNet uses two variants of attention blocks:

4.1.1 Hybrid Attention Mechanisms

- **Channel Attention Module (CEM):** Computes channel dependencies with $O(C^2)$ complexity. Given feature map $\mathbf{X} \in \mathbb{R}^{C \times H \times W}$, queries $\mathbf{Q}$, keys $\mathbf{K}$, and values $\mathbf{V}$ are generated via 1×1 convolutions followed by 3×3 depthwise convolutions. The channel attention map $\mathbf{A}_c$ is derived as:

$$\mathbf{A}_c = \text{Softmax}\left(\text{ReLU}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\alpha}\right)\right)\mathbf{V} \tag{2}$$

  where α is a learnable scaling parameter.

Table 1: Data Sources and Specifications

Data Source	**Res.**	**Role / Type**	**Bands Used**	**Description / Pipeline Context**
Landsat-8 Fire Masks	30 m	Target / Ground Truth	Active Fire Mask generated by de Almeida Pereira et al., 2021	High-resolution active binary fire mask (256×256 px).
VIIRS S-NPP	375 m	Intermediate Target / Input in Cascaded Pipeline	I4 (3.74 μm, Active Fire)	Active fire detection product (NASA FIRMS) in Fire Radiative Power (FRP) used for cascaded downscaling.
GOES-16/17 (ABI)	2000 m	Coarse Fire Input	Ch. 7 (3.9 μm), Ch. 14 (11.2 μm)	High-temporal-frequency geostationary data in Brightness Temperature (BT).
Copernicus DEM	30 m	Covariate (Static)	DEM	Digital Elevation Model to capture terrain effects on fire spread.
ESA WorldCover	10 m	Covariate (Static)	All Landcover types	Landcover maps.

Table 2: Data Pipeline Configuration Specifications

Parameter / Constant	Value	Description
VIIRS Geographic Bounds	$[-170.0, -30.0]$	Geographic longitude limits for the Americas
VIIRS Date Range	2020-01-01 to 2020-12-31	Temporal query window for VIIRS detections
Landsat Geographic Bounds	North and South America	Locations of interest for GOES detections
Landsat Date Range	Aug and Sep 2020	Temporal window for Landsat detections (de Almeida Pereira et al., 2021)
Minimum Landsat Active Pixels	10	Minimum active Landsat fire pixels required to retain a patch
GOES Temporal Guard	60 minutes	Max allowed temporal gap between target time and GOES scan
GOES Native Scale	2000 m	Native spatial grid spacing of GOES ABI channel data
Landsat Grid Scale	30 m	Spatial resolution for Landsat processing branch
VIIRS Grid Scale	375 m	Spatial resolution for VIIRS active fire detections
Landsat Dimensions	256×256 px	Target dimensions for Landsat patches
VIIRS Dimensions	20×20 px	Target spatial dimension for VIIRS rasters (7680/375)
GOES Dimensions	4×4 px	Target spatial dimension for GOES rasters (7680/2000)

- **Spatial-Channel Attention (S_CEM):** Parallel spatial attention branch operating along transposed dimensions $\mathbf{Q}_s, \mathbf{K}_s, \mathbf{V}_s \in \mathbb{R}^{B \times N_{head} \times HW \times C}$. The full spatial attention step scales quadratically with spatial resolution as $O((HW)^2)$.

4.1.2 Full-Resolution Covariate Pyramid Branch

Because full spatial attention in S_CEM imposes an $O((HW)^2)$ memory footprint that causes out-of-memory errors on full-resolution inputs, the inputs (256×256) are downsampled by factor $D = 4$ (i.e. $256/4 = 64$) before passing to the model.

To prevent losing high-resolution spatial information of the covariates, we introduce a parallel, full-resolution Covariate Pyramid Encoder. This branch processes full-resolution covariate features using Windowed Spatial-Channel Attention Blocks (Windowed-S_CEM) (Liu et al., 2021) for computational efficiency.

For a feature map partitioned into non-overlapping local windows of size $P \times P$ ($P = 8$), Windowed-S_CEM computes spatial attention independently within each window. The window is shifted sequentially with shift size $P/2$. This keeps the spatial attention complexity to $O(P^4)$ per window, independent of overall spatial input dimensions.

4.1.3 Multi-Scale Refinement and Feature Fusion

The main encoder-decoder path produces low-resolution feature representations (64×64). The final refinement stage produces the final high-resolution output. At each stage:

1. Low-resolution features are upsampled via bilinear interpolation followed by 3×3 convolutions.

2. Upsampled features are concatenated with matching scale representations extracted from the Covariate Pyramid Encoder.

3. A 3×3 fusion convolution maps concatenated features back to channel dimension.

5 Results

6 Discussions

7 Conclusions

References

Chen, H.-W., Xu, Y.-S., Hong, M.-F., Tsai, Y.-M., Kuo, H.-K., & Lee, C.-Y. (2023). Cascaded Local Implicit Transformer for Arbitrary-Scale Super-Resolution, 18257–18267. Retrieved June 7, 2026, from https://openaccess.thecvf.com/content/CVPR2023/html/Chen_Cascaded_Local_Implicit_Transformer_for_Arbitrary-Scale_Super-Resolution_CVPR_2023_paper.html

de Almeida Pereira, G. H., Fusioka, A. M., Nassu, B. T., & Minetto, R. (2021). Active fire detection in Landsat-8 imagery: A large-scale dataset and a deep-learning study. *ISPRS Journal of Photogrammetry and Remote Sensing*, *178*, 171–186. https://doi.org/10.1016/j.isprsjprs.2021.06.002

Dwyer, O., & Nam, H. W. (2026, March 31). *How wildfires and storms drove insurance losses in 2025 – in three charts*. Carbon Brief. Retrieved June 10, 2026, from https://www.carbonbrief.org/how-wildfires-and-storms-drove-insurance-losses-in-2025-in-three-charts/

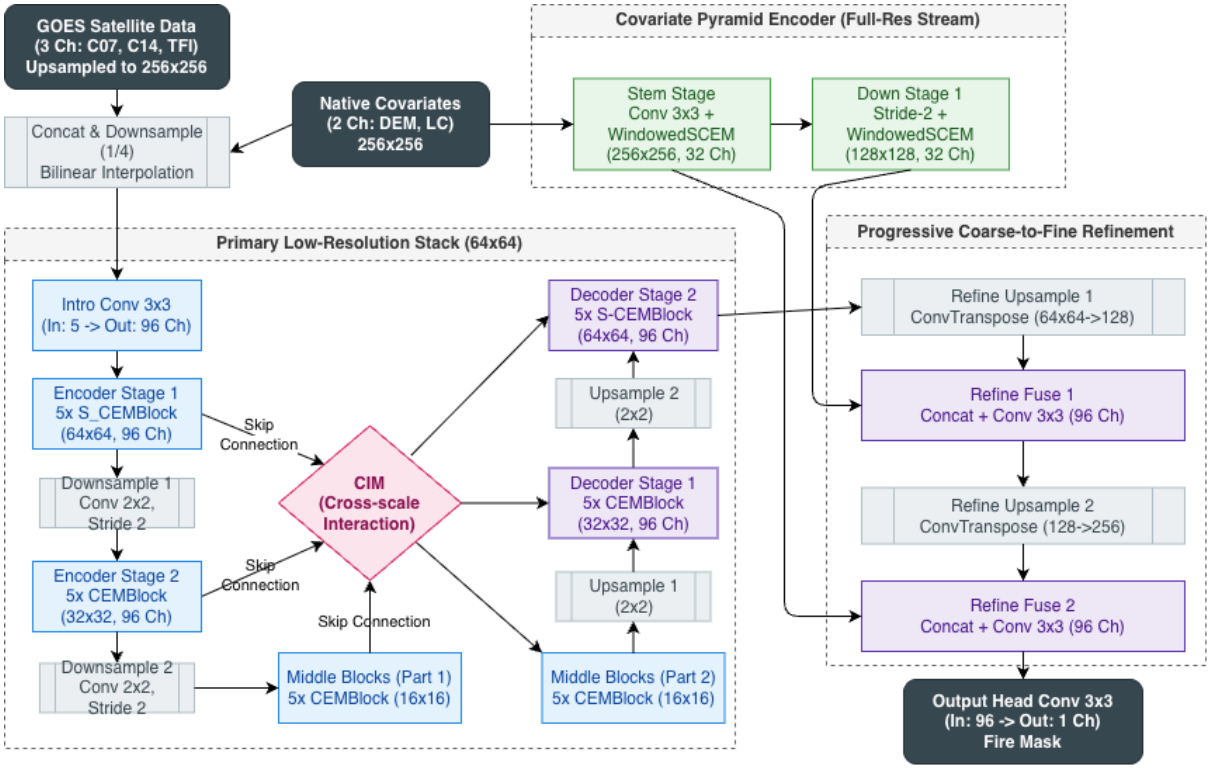


Figure 1: Architecture overview of the wildfire-adapted HAUNet. The full-resolution stream processes native covariates, which are progressively fused via refinement blocks with the low-resolution stack after CIM.

Jing, Y., Lin, L., Li, X., Li, T., & Shen, H. (2022). Cascaded Downscaling–Calibration Networks for Satellite Precipitation Estimation. *IEEE Geoscience and Remote Sensing Letters*, *19*, 1–5. https://doi.org/10.1109/LGRS.2022.3214083

Jones, M. W., Smith, A. J. P., Betts, R., Canadell, J. G., Prentice, I. C., & Quéré, C. L. (2020). Climate Change Increases the Risk of Wildfires.

Lambert, J., & Turner, G. (2024, November 14). *Investment Trends and Outcomes in the Global Carbon-Credit Market-2024 | MSCI*. Retrieved June 5, 2026, from https://www.msci.com/research-and-insights/paper/investment-trends-and-outcomes-in-the-global-carbon-credit-market-2024

Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., Lin, S., & Guo, B. (2021). Swin Transformer: Hierarchical Vision Transformer using Shifted Windows. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 9992–10002. https://doi.org/10.1109/ICCV48922.2021.00986

Qi, Y., Lou, M., Liu, Y., Li, L., Yang, Z., & Nie, W. (2026). Advancing image super-resolution techniques in remote sensing: A comprehensive survey. *ISPRS Journal of Photogrammetry and Remote Sensing*, *231*, 68–100. https://doi.org/10.1016/j.isprsjprs.2025.10.024

Sdraka, M., Papoutsis, I., Psomas, B., Vlachos, K., Ioannidis, K., Karantzalos, K., Gialampoukidis, I., & Vrochidis, S. (2022). Deep Learning for Downscaling Remote Sensing Images: Fusion and super-resolution. *IEEE Geoscience and Remote Sensing Magazine*, *10*(3), 202–255. https://doi.org/10.1109/MGRS.2022.3171836

Wang, J., Chen, M., Karaev, N., Vedaldi, A., Rupprecht, C., & Novotny, D. (2025, March 14). *VGGT: Visual Geometry Grounded Transformer*. arXiv: 2503.11651 [cs.CV]. https://doi.org/10.48550/arXiv.2503.11651

Wang, J., Wang, B., Wang, X., Zhao, Y., & Long, T. (2023). Hybrid Attention-Based U-Shaped Network for Remote Sensing Image Super-Resolution. *IEEE Transactions on Geoscience and Remote Sensing*, *61*, 1–15. https://doi.org/10.1109/TGRS.2023.3283769

Wang, Y., Yang, W., Chen, X., Wang, Y., Guo, L., Chau, L.-P., Liu, Z., Qiao, Y., Kot, A. C., & Wen, B. (2023, November 23). *SinSR: Diffusion-Based Image Super-Resolution in a Single Step*. arXiv: 2311.14760 [cs.CV]. https://doi.org/10.48550/arXiv.2311.14760

Wang, Z., Chen, J., & Hoi, S. C. H. (2020, February 8). *Deep Learning for Image Super-resolution: A Survey*. arXiv: 1902.06068 [cs.CV]. https://doi.org/10.48550/arXiv.1902.06068